

MATH7030 Advanced Numerical Methods

Chapter 1 Introduction to Numerical Computation

BY YULIANG WANG



北京师范大学 香港浸会大学 联合国际学院

BEIJING NORMAL UNIVERSITY · HONG KONG BAPTIST UNIVERSITY
UNITED INTERNATIONAL COLLEGE

Email: yuliangwang@uic.edu.cn



1 Floating-point numbers and arithmetic



- Decimal Number System (Base 10)

$$123.45 = 1 \times 10^2 + 2 \times 10^1 + 3 \times 10^0 + 4 \times 10^{-1} + 5 \times 10^{-2}$$

- Binary number system (Base 2)

$$1101_2 = 2^3 + 2^2 + 0 + 2^0 = 13$$

$$1101.11_2 = 2^3 + 2^2 + 0 + 2^0 + 2^{-1} + 2^{-2} = 13.75$$


Exercise 1.1. Repeating expansions.

Exercise 1.2. Arithmetic operations on binary numbers.

Exercise 1.3. Convert numbers from base 10 to base 2.



- Fixed-point representation has a fixed number of bits for the integral and fractional parts.
- There are three parts of the fixed-point number representation: **Sign bit**, **Integer part**, and **Fractional part**.



- Sign bit: the negative number has a sign bit 1, while a positive number has a bit 0.
- Integer Part: its length depends on the size of the computer word; for an 8-bit word, the integral part is 4 bits.
- Fractional part: for an 8-bit word, the fractional part is 3 bits.


Example 1.1 Write 4.5 in 8-bit fixed-point representation.



- A floating-point representation has three parts: **Sign bit**, **Exponent Part**, and **Mantissa**:

$$\pm m \times b^E$$

- m : mantissa/significand
- b : base
- E : exponent



- A **single-precision** word consists of 32 bits: 1 bit for the sign (0 for +, 1 for −), 8 bits for the exponent, and 23 bits for the mantissa.
- The mantissa is always **normalized** so that $1 \leq m < 2$, e.g. $10 = 1010_2 = 1.010 \times 2^3$. So the first bit of the mantissa is always 1.



- A number that can be stored exactly using this scheme is called a **floating-point number**, e.g. $1010_2 = 1.010 \times 2^3$ is a floating-point number:

0	$E = 3$	10100...0
---	---------	-----------

The exact format of E is discussed later.

- Since the first bit in m is always 1, one can omit to obtain the **hidden-bit representation**:

0	$E = 3$	0100...0
---	---------	----------

- $0.1 = 0.000\overline{1100}_2 = 1.100\overline{1100}_2 \times 2^{-4}$ is not a floating-point number since it must be rounded in order to be stored. In single-precision, $1.1001100110011001100\backslash 1100\overline{1100}$ is approximated by either


0	$E = -4$	10011001100110011001100
---	----------	-------------------------

or

0	$E = -4$	10011001100110011001101
---	----------	-------------------------



- The gap between 1 and the next larger floating-point number is called the **machine precision** and is often denoted by ϵ .
- In single precision, the next floating-point number after 1 is $1 + 2^{-23}$. Thus, for single precision, we have $\epsilon = 2^{-23} \approx 1.2 \times 10^{-7}$.
- The default precision in MATLAB is **double precision**, where a word consists of 64 bits: 1 for the sign, 11 for the exponent, and 52 for the significand. Hence in double precision the next larger floating-point number after 1 is $1 + 2^{-52}$, so that the machine precision is $2^{-52} \approx 2.2 \times 10^{-16}$.
- Consider a toy system in which only significands of the form $1.b_1b_2$ can be represented and only exponents 0, 1, and -1 can be stored. What are the numbers that can be represented in this system?





- the machine precision ϵ is 0.25.
- In general, the gaps between representable numbers become larger as we move away from the origin. This is acceptable since the relative gaps remains of reasonable size.
- However, the gap between 0 and the smallest positive number is much larger than the gap between the smallest and next smallest positive number.
- This is the case with single-and double-precision floating-point numbers as well.
- This gap can be filled in using **subnormal** numbers.



- A special representation is needed for 0,
- and also for $\pm\infty$ (the result of dividing a nonzero number by 0),
- and also for NaN (Not a Number; e.g., $0/0$).
- These are done with special bits in the exponent field.
- Special bits in the exponent field are also used to signal subnormal numbers.



If exponent field is	Then number is	Type of number:
00000000000	$\pm(0.b_1 \dots b_{52})_2 \times 2^{-1022}$	0 or subnormal
00000000001 = 1_{10} 00000000010 = 2_{10}	$\pm(1.b_1 \dots b_{52})_2 \times 2^{-1022}$ $\pm(1.b_1 \dots b_{52})_2 \times 2^{-1021}$	Normalized number
\vdots 01111111111 = 1023_{10} \vdots	\vdots $\pm(1.b_1 \dots b_{52})_2 \times 2^0$ \vdots	Exponent field is (actual exponent) + 1023
11111111110 = 2046_{10}	$\pm(1.b_1 \dots b_{52})_2 \times 2^{1023}$	
11111111111	$\pm\infty$ if $b_1 = \dots = b_{52} = 0$, NaN otherwise	Exception



- The smallest positive normalized floatingpoint number that can be stored is $1.0_2 \times 2^{-1022} \approx 2.2 \times 10^{-308}$, while the largest is $1.1\dots1_2 \times 2^{1023} \approx 1.8 \times 10^{308}$.
- The exponent field for normalized floatingpoint numbers represents the actual exponent plus 1023. We can represent exponents between -1022 and $+1023$.
- The two special exponent field bit patterns are all 0 s and all 1 s.
- An exponent field consisting of all 0s signals either 0 or a subnormal number.
- The smallest positive subnormal number that can be represented is $2^{-52} \times 2^{-1022} = 2^{-1074}$.
- The number 0 is represented by an exponent field consisting of all 0s and a significand field of all 0s.
- An exponent field consisting of all 1s signals an exception. If all bits in the significand are 0, then it is $\pm\infty$. Otherwise it represents NaN.



If x is a real number that cannot be stored exactly, then it is replaced by a nearby floating-point number according to one of the following rules:

1. **Round down:** $\text{round}(x)$ is the largest floating-point number that is less than or equal to x .
2. **Round up:** $\text{round}(x)$ is the smallest floating-point number that is greater than or equal to x .
3. **Round towards 0:** $\text{round}(x)$ is either the round down or round up of x , whichever lies between 0 and x .
4. **Round to nearest:** $\text{round}(x)$ is either the round down or round up of x , whichever is closer. In case of a tie (x falls in the middle), it is the one whose least significant (rightmost) bit is 0 (ties to even).

The default is round to the nearest.



Exercise 1.4. Using double precision, find the floating-point representation of $0.1 = 1.100\overline{1100}_2 \times 2^{-4}$ under the four rounding modes.

- The absolute rounding error associated with a number x is defined as $|\text{round}(x) - x|$.
- If $x = \pm m \times 2^E$ (excluding subnormal numbers), then the absolute rounding error is always $< \epsilon \times 2^E$ for any rounding mode, and $\leq \frac{\epsilon}{2} \times 2^E$ for round to nearest, where ϵ is the machine precision.
- The relative rounding error associated with a number x is defined as $|\text{round}(x) - x| / |x|$, which is always $< \epsilon$ for any rounding mode, and $\leq \frac{\epsilon}{2}$ for round to nearest. In other words, we can write

$$\text{round}(x) = x(1 + \delta), \quad \text{where } |\delta| < \epsilon \quad (\text{or } \leq \frac{\epsilon}{2} \text{ for round to nearest}).$$



The IEEE standard requires that the result of an operation (addition, subtraction, multiplication, or division) on two floating-point numbers must be the **correctly rounded** value of the exact result.

It means that if a and b are floatingpoint numbers and \oplus , \ominus , \otimes , and \oslash represent floating-point addition, subtraction, multiplication, and division, then we will have

$$a \oplus b = \text{round}(a + b) = (a + b)(1 + \delta_1),$$

$$a \ominus b = \text{round}(a - b) = (a - b)(1 + \delta_2),$$

$$a \otimes b = \text{round}(ab) = (ab)(1 + \delta_3),$$

$$a \oslash b = \text{round}(a/b) = (a/b)(1 + \delta_4),$$

where $|\delta_i| < \epsilon$ (or $\leq \epsilon/2$ for round to nearest), $i = 1, \dots, 4$.



- **Overflow** occurs when the true result of an operation is greater than the largest floating-point number ($1.1\dots1_2 \times 2^{1023} \approx 1.8 \times 10^{308}$ for double precision).
 - Using round up or round to nearest, the result is set to ∞ ;
 - using round down or round towards 0, it is set to the largest floating-point number.
- **Underflow** occurs when the true result is less than the smallest floating-point number. The result is stored as a subnormal number if it is in the range of the subnormal numbers, and otherwise it is set to 0.
- Other operations that produce ∞ : for any $x > 0$ or $x = \infty$, $\frac{x}{0}$, $x + \infty$, $x \times \infty$ all produce ∞ . Similar rules apply to $-\infty$.
- Some operations to produce NaN (incomplete list):

$$\infty - \infty, \quad -\infty + \infty, \quad 0 \times \infty, \quad 0 \div 0, \quad \infty \div \infty, \quad \sqrt{-1}$$



2 Conditioning of problems



Types of errors in scientific computing:

1. physical model \longrightarrow mathematical model (simplifications etc.)
2. mathematical model \longrightarrow numerical model (discretization, truncation etc.)
3. numerical model \longrightarrow numerical algorithm, implemented on a computer (rounding errors etc.)
4. errors in input data (measurement, inverse problems)

Given the true value y and the computed value \hat{y} ,

- $|\hat{y} - y|$ is called the **absolute error**
- $\frac{|\hat{y} - y|}{|y|}$ is called the **relative error**



The **conditioning** of a problem measures how sensitive the output is to small changes in the input.

Let $f: \mathbb{R} \rightarrow \mathbb{R}$. Consider the problem: given $x \in \mathbb{R}$, find $y = f(x)$.

If $\hat{x} \neq x$ (but close to x), and $\hat{y} = f(\hat{x})$. How close is \hat{y} to y ?

If

$$|\hat{y} - y| \approx C(x)|\hat{x} - x|,$$

then we might call $C(x)$ the **absolute condition number** of the function f at x .

If

$$\left| \frac{\hat{y} - y}{y} \right| \approx \kappa(x) \left| \frac{\hat{x} - x}{x} \right|,$$

then $\kappa(x)$ might be called the **relative condition number** of f at x .



If $f'(x)$ exists, we can take

$$C(x) = |f'(x)|, \quad \kappa(x) = \left| \frac{xf'(x)}{f(x)} \right|.$$

Example 2.1 Find $C(x)$ and $\kappa(x)$ for $f(x) = 2x$.

Example 2.2 Find $C(x)$ and $\kappa(x)$ for $f(x) = \sqrt{x}$.

Example 2.3 Find the $C(x)$ and $\kappa(x)$ for solving the linear system $Ax = b$, assuming A is nonsingular and fixed.



An algorithm that achieves the level of accuracy defined by the conditioning of the problem is called **stable**, while one that gets unnecessarily inaccurate results called **unstable**.

Example 2.4 (Computing sums) If x and y are two real numbers and they are rounded to floating-point numbers and their sum is computed on a machine with machine precision ϵ .

$$\text{fl}(x + y) \equiv \text{round}(x) \oplus \text{round}(y) = (x(1 + \delta_1) + y(1 + \delta_2))(1 + \delta_3), \quad |\delta_i| \leq \epsilon.$$

where $\text{fl}(\cdot)$ denotes the floating-point result.



Forward error analysis. How much does the computed value differ from the exact solution?

- absolute error: $|\text{fl}(x + y) - (x + y)| \leq (|x| + |y|)(2\epsilon + \epsilon^2)$: small
- relative error: $\left| \frac{\text{fl}(x + y) - (x + y)}{x + y} \right| \leq \frac{(|x| + |y|)(2\epsilon + \epsilon^2)}{|x + y|}$: can be large if $y \approx -x$, even if $\delta_3 = 0$, so the problem is ill-conditioned in this case.

Backward error analysis. Here one tries to show that the computed value is the exact solution to a nearby problem.

$$\text{fl}(x + y) = x(1 + \delta_1)(1 + \delta_3) + y(1 + \delta_2)(1 + \delta_3)$$

the computed value is the exact sum of two numbers that differ from x and y by relative amounts no greater than $2\epsilon + \epsilon^2$.



Example 2.5 Compute $\exp(x)$ using the Taylor series expansion

$$e^x = 1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + \cdots$$

```
function newsum = my_exp(x)
oldsum = 0;
newsum = 1;
term = 1;
n = 0;
while newsum ~= oldsum % Iterate until next term is negligible
    n = n + 1;
    term = term * x/n; % x^n/n! = (x^{n-1}/(n-1)!) * x/n
    oldsum = newsum;
    newsum = newsum + term;
end
end
```

high relative error for large negative values of x



The problem itself is well-conditioned (check).

But the algorithm is unstable.

Exercise 2.1. Modify the code to make it stable for all x .

Numerical differentiation. Forward difference

$$f(x+h) = f(x) + hf'(x) + \frac{h^2}{2}f''(\xi), \quad \xi \in [x, x+b]$$

$$f'(x) = \frac{f(x+h) - f(x)}{h} - \frac{h}{2}f''(\xi).$$

The term $-\frac{h}{2}f''(\xi) = O(h)$ is referred to as the **truncation error** or discretization error.



But the first term may also contain error in numerical computation.

$$\frac{f(x+h)(1+\delta_1) - f(x)(1+\delta_2)}{h} = \frac{f(x+h) - f(x)}{h} + \frac{\delta_1 f(x+h) - \delta_2 f(x)}{h}, \quad |\delta_i| < \epsilon$$

$$\left| \frac{\delta_1 f(x+h) - \delta_2 f(x)}{h} \right| \lesssim \frac{2\epsilon |f(x)|}{h}.$$

Example 2.6 Suppose $f(x) = \sin x$. What's the best accuracy by approximating $f'(\pi/4)$ using the forward difference?

Higher order finite difference (central difference, etc.)